

- When conducting the work of Lab 11, we conducted the test that uses the Central Limit Theorem even though the sample size was “small” (i.e., $n < 30$). It turns out, that how “far off” the t -test is can be computed using a first-order Edgeworth approximation for the error. Below, we will do this for the the further observations.

- Boos and Hughes-Oliver (2000) note that

$$P(T \leq t) \approx F_Z(t) + \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

where $f_Z(\cdot)$ and $F_Z(\cdot)$ are the Gaussian PDF and CDF and skew is the skewness of the data. What is the potential error in the computation of the p -value when testing $H_0 : \mu_X = 0; H_a : \mu_X < 0$ using the zebra finch further data?

```
skewn = skewness(data$farther)

tstat = t.test(data$farther, alternative = "less")["statistic"]["t"]

pdf = dnorm(tstat)

(error = (pdf)*(1/5)*skewn*(2*(tstat)^2 + 1)*(1/6))

## [1] -1.226006e-13
```

$$\text{Error} = -1.226006e - 13$$

- Compute the error for t statistics from -10 to 10 and plot a line that shows the error across t . Continue to use the skewness and the sample size for the zebra finch further data.

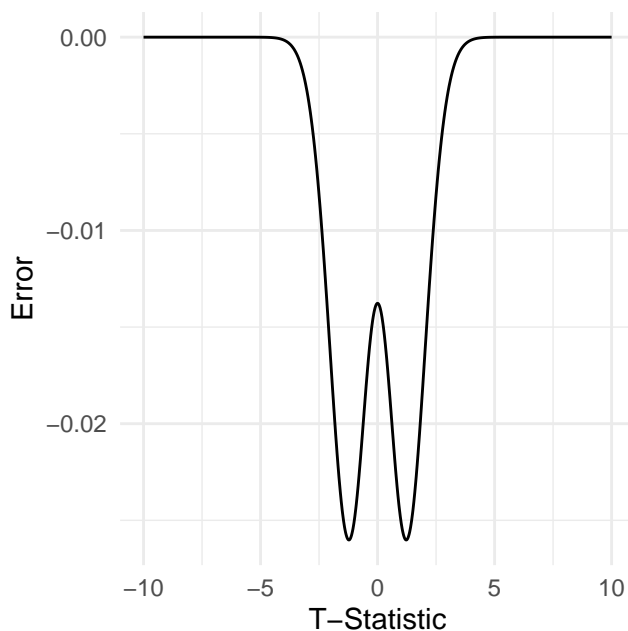


Figure 1: Error Graph Based on t

- Suppose we wanted to have a tail probability within 10% of the desired $\alpha = 0.05$. Recall we did a left-tailed test using the further data. How large of a sample size would we need? That is, we

need to solve the error formula equal to 10% of the desired left-tail probability:

$$0.10\alpha \stackrel{set}{=} \underbrace{\frac{\text{skew}}{\sqrt{n}} \frac{(2t^2 + 1)}{6} f_Z(t)}_{\text{error}},$$

which yields

$$n = \left(\frac{\text{skew}}{6(0.10\alpha)} (2t^2 + 1) f_Z(t) \right)^2.$$

```
alpha = 0.05

(size = ((skewness(data$farther)/(6*0.1*alpha))*(2*qnrm(alpha)^2 + 1)*dnrm(qnrm(alpha)))^2)

## [1] 520.8876
```

$$n = 520.8876$$

2. Complete the following steps to revisit the analyses from lab 11 using the bootstrap procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform resampling to approximate the sampling distribution of the T statistic:

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}},$$

where \bar{x}_r is the mean computed on the r^{th} resample and s is the sample standard deviation from the original samples. At the end, create an object called `resamples.null.closer`, for example, and store the resamples shifted to ensure they are consistent with the null hypotheses at the average (i.e., here ensure the shifted resamples are 0 on average, corresponding to $t = 0$, for each case).

```
#FARTHER
farther.resampled.t = c()
farther.resamples.shifted = c()
farther.resampled.p = c()
farther.resampled.mean = c()

for(i in 1:10000){
  curr.sample = sample(data$farther, size = length(data$farther), replace = T)
  farther.resampled.mean[i] = mean(curr.sample)
  farther.resampled.t[i] = 5*farther.resampled.mean[i]/sd(data$farther)
}

for(i in 1:10000){
  farther.resamples.shifted[i] = farther.resampled.t[i] - mean(farther.resampled.t)
}

#CLOSER
closer.resampled.t = c()
closer.resamples.shifted = c()
closer.resampled.p = c()
closer.resampled.mean = c()

for(i in 1:10000){
  curr.sample = sample(data$closer, size = length(data$closer), replace = T)
  closer.resampled.mean[i] = mean(curr.sample)
  closer.resampled.t[i] = 5*closer.resampled.mean[i]/sd(data$closer)
}

for(i in 1:10000){
  closer.resamples.shifted[i] = closer.resampled.t[i] - mean(closer.resampled.t)
}

#DIFFERENCE
dif.resampled.t = c()
dif.resamples.shifted = c()
dif.resampled.p = c()
dif.resampled.mean = c()

for(i in 1:10000){
```

```

curr.sample = sample(data$difference, size = length(data$difference), replace = T)
dif.resampled.mean[i] = mean(curr.sample)
dif.resampled.t[i] = 5*dif.resampled.mean[i]/sd(data$difference)

}
for(i in 1:10000){
dif.resamples.shifted[i] = dif.resampled.t[i] - mean(dif.resampled.t)
}

```

- (b) Compute the bootstrap p -value for each test using the shifted resamples. How do these compare to the t -test p -values?

```

#FARTHER
for(i in 1:10000){
farther.resampled.p[i] = length(which(farther.resamples.shifted <= farther.resampled.t[i]))/length(farther.resamples.shifted)
}
mean(farther.resampled.p)

## [1] 0

#CLOSER
for(i in 1:10000){
closer.resampled.p[i] = length(which(closer.resamples.shifted >= closer.resampled.t[i]))/length(closer.resamples.shifted)
}
mean(closer.resampled.p)

## [1] 0

#DIFFERENCE
for(i in 1:10000){
dif.resampled.p[i] = length(which(dif.resamples.shifted >= abs(dif.resampled.t[i]) | dif.resamples.shifted <= -abs(dif.resampled.t[i])))/length(dif.resamples.shifted)
}
mean(dif.resampled.p)

## [1] 0

```

$$p_f = p_c = p_d = 0$$

- (c) What is the 5th percentile of the shifted resamples under the null hypothesis? Note this value approximates $t_{0.05, n-1}$. Compare these values in each case.

```

(farther.quantile = quantile(farther.resamples.shifted, 0.05))

##      5%
## -1.7118

(closer.quantile = quantile(closer.resamples.shifted, 0.05))

##      5%
## -1.56891

(dif.quantile = quantile(dif.resamples.shifted, 0.05))

##      5%
## -1.568147

```

$$t_f = -1.672723$$

$$t_c = -1.612601$$

$$t_d = -1.569826$$

- (d) Compute the bootstrap confidence intervals using the resamples. How do these compare to the t -test confidence intervals?

```

quantile(farther.resampled.mean, c(0.025, 0.975))

##      2.5%      97.5%
## -0.2566202 -0.1552488

quantile(closer.resampled.mean, c(0.025, 0.975))

##      2.5%      97.5%
##  0.1209154  0.1926862

```

```
quantile(dif.resampled.mean, c(0.025,0.975))

##      2.5%      97.5%
## 0.2811266 0.4418409
```

$$CI_f = (-0.2572578, -0.1554500)$$

$$CI_c = (0.1208817, 0.1918892)$$

$$CI_d = (0.2791661, 0.4410959)$$

3. Complete the following steps to revisit the analyses from lab 11 using the randomization procedure.

- (a) Now, consider the zebra finch data. We do not know the generating distributions for the closer, further, and difference data, so perform the randomization procedure

```
#FARTHER
farther.xbars = c()

for(i in 1:10000){
  curr.rand = data$farther*sample(c(-1,1),length(data$farther),replace = T)
  farther.xbars[i] = mean(curr.rand)
}

#CLOSER
closer.xbars = c()

for(i in 1:10000){
  curr.rand = data$closer*sample(c(-1,1),length(data$closer),replace = T)
  closer.xbars[i] = mean(curr.rand)
}

#DIFFERENCE
dif.xbars = c()

for(i in 1:10000){
  curr.rand = data$difference*sample(c(-1,1),length(data$difference),replace = T)
  dif.xbars[i] = mean(curr.rand)
}
```

- (b) Compute the randomization test p -value for each test.

```
#FARTHER
delta = abs(mean(data$farther))
low = -delta
high = delta
mean(farther.xbars <= low)

## [1] 0

#CLOSER
delta = abs(mean(data$closer))
low = -delta
high = delta
mean(closer.xbars >= high)

## [1] 0

#DIFFERENCE
delta = abs(mean(data$difference))
low = -delta
high = delta
mean(dif.xbars >= high) + mean(dif.xbars <= low)

## [1] 0
```

$$p_f = p_c = p_d = 0$$

- (c) Compute the randomization confidence interval by iterating over values of μ_0 .
Hint: You can “search” for the lower bound from Q_1 and subtracting by 0.0001, and the upper bound using Q_3 and increasing by 0.0001. You will continue until you find the first value for which the two-sided p -value is greater than or equal to 0.05.

```

#FARTHER
farther.xbars = c()
mu0.iterate = 0.001
mu0 = mean(data$farther)
ci.data <- data$farther
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    farther.xbars[i] = mean(curr.rand)
  }
  farther.xbars <- farther.xbars + mu0

  delta = abs(mean(data$farther) - mu0)
  low = mu0 - delta
  high = mu0 + delta
  p.val = mean(farther.xbars <= low) +
    mean(farther.xbars >= high)

  if(p.val < 0.05){
    break
    mu0 <- mu0 + mu0.iterate
  }else{
    mu0 <- mu0 - mu0.iterate
  }
}
farther.rand.lbound = mu0

mu0.iterate = 0.001
mu0 = mean(data$farther)
ci.data <- data$farther
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    farther.xbars[i] = mean(curr.rand)
  }
  farther.xbars <- farther.xbars + mu0

  delta = abs(mean(data$farther) - mu0)
  low = mu0 - delta
  high = mu0 + delta
  p.val = mean(farther.xbars <= low) +
    mean(farther.xbars >= high)

  if(p.val < 0.05){
    mu0 = mu0 - mu0.iterate
    break
  }else{
    mu0 <- mu0 + mu0.iterate
  }
}
farther.rand.ubound = mu0

farther.rand.lbound

## [1] -0.2557244

farther.rand.ubound

## [1] -0.1527244

#CLOSER
mu0.iterate = 0.001
mu0 = mean(data$farther)
ci.data <- data$farther
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    farther.xbars[i] = mean(curr.rand)
  }
  farther.xbars <- farther.xbars + mu0

  delta = abs(mean(data$farther) - mu0)
  low = mu0 - delta
  high = mu0 + delta
  p.val = mean(farther.xbars <= low) +

```

```

    mean(farther.xbars >= high)

if(p.val < 0.05){
  mu0 = mu0 - mu0.iterate
  break
}else{
  mu0 <- mu0 + mu0.iterate
}
}
farther.rand.ubound = mu0

farther.rand.lbound

## [1] -0.2557244

farther.rand.ubound

## [1] -0.1517244

#Closer#####

closer.xbars = c()

for(i in 1:10000){
  curr.rand = data$closer*sample(c(-1,1),length(data$closer),replace = T)
  closer.xbars[i] = mean(curr.rand)
}

delta = abs(mean(data$closer))
low = -delta
high = delta
mean(closer.xbars >= high)

## [1] 0

closer.xbars = c()
mu0.iterate = 0.001
mu0 = mean(data$closer)
ci.data <- data$closer
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    closer.xbars[i] = mean(curr.rand)
  }
  closer.xbars <- closer.xbars + mu0

  delta = abs(mean(data$closer) - mu0)
  low = mu0 - delta
  high = mu0 + delta
  p.val = mean(closer.xbars <= low) +
    mean(closer.xbars >= high)

  if(p.val < 0.05){
    break
    mu0 <- mu0 + mu0.iterate
  }else{
    mu0 <- mu0 - mu0.iterate
  }
}
closer.rand.lbound = mu0

mu0.iterate = 0.001
mu0 = mean(data$closer)
ci.data <- data$closer
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    closer.xbars[i] = mean(curr.rand)
  }
  closer.xbars <- closer.xbars + mu0

  delta = abs(mean(data$closer) - mu0)
  low = mu0 - delta
  high = mu0 + delta
  p.val = mean(closer.xbars <= low) +
    mean(closer.xbars >= high)

```

```

if(p.val < 0.05){
  mu0 = mu0 - mu0.iterate
  break
}else{
  mu0 <- mu0 + mu0.iterate
}
}
closer.rand.ubound = mu0

closer.rand.lbound

## [1] 0.1172231

closer.rand.ubound

## [1] 0.1952231

#DIFFERENCE
dif.xbars = c()
mu0.iterate = 0.001
mu0 = mean(data$difference)
ci.data <- data$difference
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    dif.xbars[i] = mean(curr.rand)
  }
  dif.xbars <- dif.xbars + mu0

delta = abs(mean(data$difference) - mu0)
low = mu0 - delta
high = mu0 + delta
p.val = mean(dif.xbars <= low) +
  mean(dif.xbars >= high)

if(p.val < 0.05){
  break
  mu0 <- mu0 + mu0.iterate
}else{
  mu0 <- mu0 - mu0.iterate
}
}
dif.rand.lbound = mu0

mu0.iterate = 0.001
mu0 = mean(data$difference)
ci.data <- data$difference
repeat{
  curr.shifted.dat <- ci.data - mu0
  for(i in 1:1000){
    curr.rand = curr.shifted.dat*sample(c(-1,1),length(curr.shifted.dat),replace = T)
    dif.xbars[i] = mean(curr.rand)
  }
  dif.xbars <- dif.xbars + mu0

delta = abs(mean(data$difference) - mu0)
low = mu0 - delta
high = mu0 + delta
p.val = mean(dif.xbars <= low) +
  mean(dif.xbars >= high)

if(p.val < 0.05){
  mu0 = mu0 - mu0.iterate
  break
}else{
  mu0 <- mu0 + mu0.iterate
}
}
dif.rand.ubound = mu0

dif.rand.lbound

## [1] 0.2729475

dif.rand.ubound

## [1] 0.4419475

```

$$CI_f = (-0.2577244, -0.1507244)$$

$$CI_c = (0.1172231, 0.1942231)$$

$$CI_d = (0.2719475, 0.4459475)$$

4. **Optional Challenge:** In this lab, you performed resampling to approximate the sampling distribution of the T statistic using

$$T = \frac{\bar{x}_r - 0}{s/\sqrt{n}}.$$

I'm curious whether it is better/worse/similar if we computed the statistics using the sample standard deviation of the resamples (s_r), instead of the original sample (s)

$$T = \frac{\bar{x}_r - 0}{s_r/\sqrt{n}}.$$

- (a) Perform a simulation study to evaluate the Type I error for conducting this hypothesis test both ways.
- (b) Using the same test case(s) as part (a), compute bootstrap confidence intervals and assess their coverage – how often do we ‘capture’ the parameter of interest?

References

Boos, D. D. and Hughes-Oliver, J. M. (2000). How large does n have to be for z and t intervals? *The American Statistician*, 54(2):121–128.